

SYSC4907 PROJECT:
SENSOR-BASED ACCESS CONTROL SYSTEM

By
Craig Shorrocks, Jessica Morris, Richard Perryman
March 2017

A Fourth Year Project Report
submitted to the Dept. of Systems & Computer Engineering
in partial fulfillment of the requirements
for the degree of
Bachelors of Engineering

© Copyright 2017
by Craig Shorrocks, Jessica Morris, Richard Perryman , Ottawa, Canada

Abstract

This report tells you all you need to know about something.

Acknowledgements

I would like to thank my supervisor, anyone who paid me money, gave me equipment, etc.

Contents

Abstract	ii
Acknowledgements	iii
List of Figures	vii
List of Tables	viii
List of Abbreviations	ix
1 Introduction	1
2 The Engineering Project	2
2.1 Health and Safety	2
2.2 Engineering Professionalism	3
2.3 Project Management	3
2.4 Individual Contributions	3
2.4.1 Project Contributions	3
2.4.2 Report Contributions	3
3 Technical Background	4
3.1 NFC	4
3.2 Cloud Computing	4
3.2.1 AWS	4
3.3 Security	4

3.4	Single-Board Computers and Microcontrollers	4
3.4.1	Raspberry Pi	4
3.4.2	Arduino	5
4	Business Use Cases	6
4.1	Online Order Secure Pickup	6
4.2	Central Mail Package Pickup	9
4.3	Long Term Storage	9
4.4	Service Provider	9
5	Problem Analysis and System Design	10
5.1	Overall System Analysis	10
5.2	NFC	11
5.2.1	Card Readers	11
5.2.2	Mobile Devices	11
5.2.3	Protocol	12
5.3	Android	13
5.3.1	User Interface	13
5.3.2	Notifications	14
5.3.3	NFC Handler	14
5.4	Lock Hardware	15
5.4.1	Core Circuit	15
5.4.2	Core Software	15
5.4.3	Authentication Modules	15
5.5	Cloud	17
5.5.1	Amazon Web Servies	17
5.5.2	Representational State Transer	18
5.5.3	JavaScript Object Notation	18
5.5.4	Node.js	19
5.5.5	MySQL	19
5.5.6	Relational Database Theory	19
5.5.7	Authentication for Server API	20

5.5.8	Storing Authenticator Data Securely	20
5.5.9	Sending Data Securely Over the Network	20
5.6	Lock Demonstration	21
6	System Implementation	22
6.1	Android	22
6.1.1	Server Interfacing	22
6.1.2	HMAC	23
6.1.3	Entity Representation	23
6.2	Hardware	24
6.2.1	Lock Control Circuit	24
6.2.2	Lock Control Software	27
6.2.3	Authentication Modules	27
6.3	Cloud	27
7	Testing and Bug Fixes	28
8	Conclusions	29
A	Sample Appendix	30
	Bibliography	30

List of Figures

1	Online Order Secure Pickup Use Case Diagram	8
2	Hardware Block Diagram	16
3	Relay Circuit Diagram [5]	26

List of Tables

List of Abbreviations

AID	Application IDentifier
APDU	Application Protocol Data Unit
API	Application Programming Interface
HMAC	Keyed-Hash Message Authentication Code
IO	Input/Output
JSON	JavaScript Object Notation
MVC	Model-View-Controller
NFC	Near Field Communication
RFC	Request for Comments

Chapter 1

Introduction

Give an introduction to your project. This might include:

- Motivation for your project
- Problem you are trying to solve
- Scope of your project
- Organization of your report

You should tune this appropriately for what best suits your project.

Chapter 2

The Engineering Project

2.1 Health and Safety

The main health and safety concern for this project was ensuring the safe preparation of circuits for the hardware component. The chance of bodily harm due to electric shock was possible, as the highest-risk component of the project drew 650 mA of current at 12V direct current. Since 650 mA is above the 500 mA threshold for possible heart fibrillation [1], the following standards were enforced to mitigate safety risks when handling the electrical components:

- The workspace and components were kept dry at all times
- Participants ensured they were sufficiently grounded before handling components
- Ground pins were connected first, to ensure any charged components would discharge
- Components' wiring was not altered while the system was powered on

2.2 Engineering Professionalism

Using their course experience of ECOR 4995 Professional Practice, students should demonstrate how their professional responsibilities were met by the goals of their project and/or during the performance of their project.

2.3 Project Management

One of the goals of the engineering project is real experience in working on a long-term team project. Students should explain what project management techniques or processes were used to coordinate, manage and perform their project.

2.4 Individual Contributions

This section should carefully itemize the individual contributions of each team member. Project contributions should identify which components of work were done by each individual. Report contributions should list the author of each major section of this report.

2.4.1 Project Contributions

Give the individual contributions of the each team member towards the project.

2.4.2 Report Contributions

Give the individual contributions of the each team member towards writing the final report.

Chapter 3

Technical Background

3.1 NFC

3.2 Cloud Computing

3.2.1 AWS

3.3 Security

3.4 Single-Board Computers and Microcontrollers

Insert blurb about what is a single-board thingamabob

3.4.1 Raspberry Pi

The Raspberry Pi is a credit-card-sized, single-board computer developed by the Raspberry Pi Foundation. In its initial design, it was intended to be a learning tool to introduce computer science to students at the pre-university level. However, its accessible price gave it increased appeal among hobby electronics enthusiasts, and it is now often used for robotics projects, media streaming, IoT projects, and more. Today, the Raspberry Pi has sold over ten million units [CITATION], with models

ranging from the 5 USD Raspberry Pi Zero, to the 35 USD Raspberry Pi 3 Model B.

3.4.2 Arduino

Chapter 4

Business Use Cases

Design for our system began with considering use cases that could be handled by our system. This way, we could be sure to only design features that would prove useful to some potential users, while also making sure that we don't miss any features that are necessary. These use cases would also provide benefit to demonstrating our system, in that stepping through a sample use case would show the value of the system to people without them requiring a strong technical background.

We primarily focused on the four use cases that follow. These were outlined in our initial proposal and have remained largely unchanged since then. They draw inspiration from situations that we either encountered ourselves or found through common knowledge and research. Online shopping seemed like a ripe market for technology like the SBACS, and mail delivery seemed that it followed a fairly similar system. Long term storage facilities also seemed as though they would benefit from abstracting away the idea of keys. Once we had considered these use cases, we found that several of the problems that they shared could be solved by a company which provided SBACS as a service.

4.1 Online Order Secure Pickup

With online shopping becoming more popular [], the hassle of being physically present to shop in stores could be eliminated by allowing customers to make an order and

payment online, then pick up their products at the store using a SBACS-secured storage container. After a customer's order has been confirmed by the store to be packed and ready for pick-up, a store employee could create a link between their customer and the storage container that holds that user's purchases. When the customer arrives at the store to pick up their items, they can gain access to the storage container assigned to them by using the authentication methods that they had previously provided as the components of their identity. Once the customer has provided these authentication units correctly, the link between their identity and the storage container is broken. This way, the user won't be able to inappropriately access the storage container again, and a new link could be made for a different customer.

A strong advantage of this system over existing ones, such as Walmart's Grab & Go, is that it creates a digital fingerprint using a combination of authentication methods supported by the SBACS, which is more difficult to break via brute force than a 6-digit PIN [1]. In addition, with the entire process of shopping being expedited, customers will spend less time in stores and parking lots, leading to less congestion due to vehicle traffic. This system will also grant customers with mobility or vision impairments increased personal autonomy, as they will be able to retrieve their goods with less aid from store employees than traditional shopping allows. Customers with impairments can perform their shopping online using their accessibility software, such as text-to-speech, instead of calling the store to request that an employee retrieve their goods. When the customer receives confirmation of their order and authentication tokens, they can access their assigned SBACS-secured locker in the store themselves, using their usual aids instead of occupying a store employee's time.

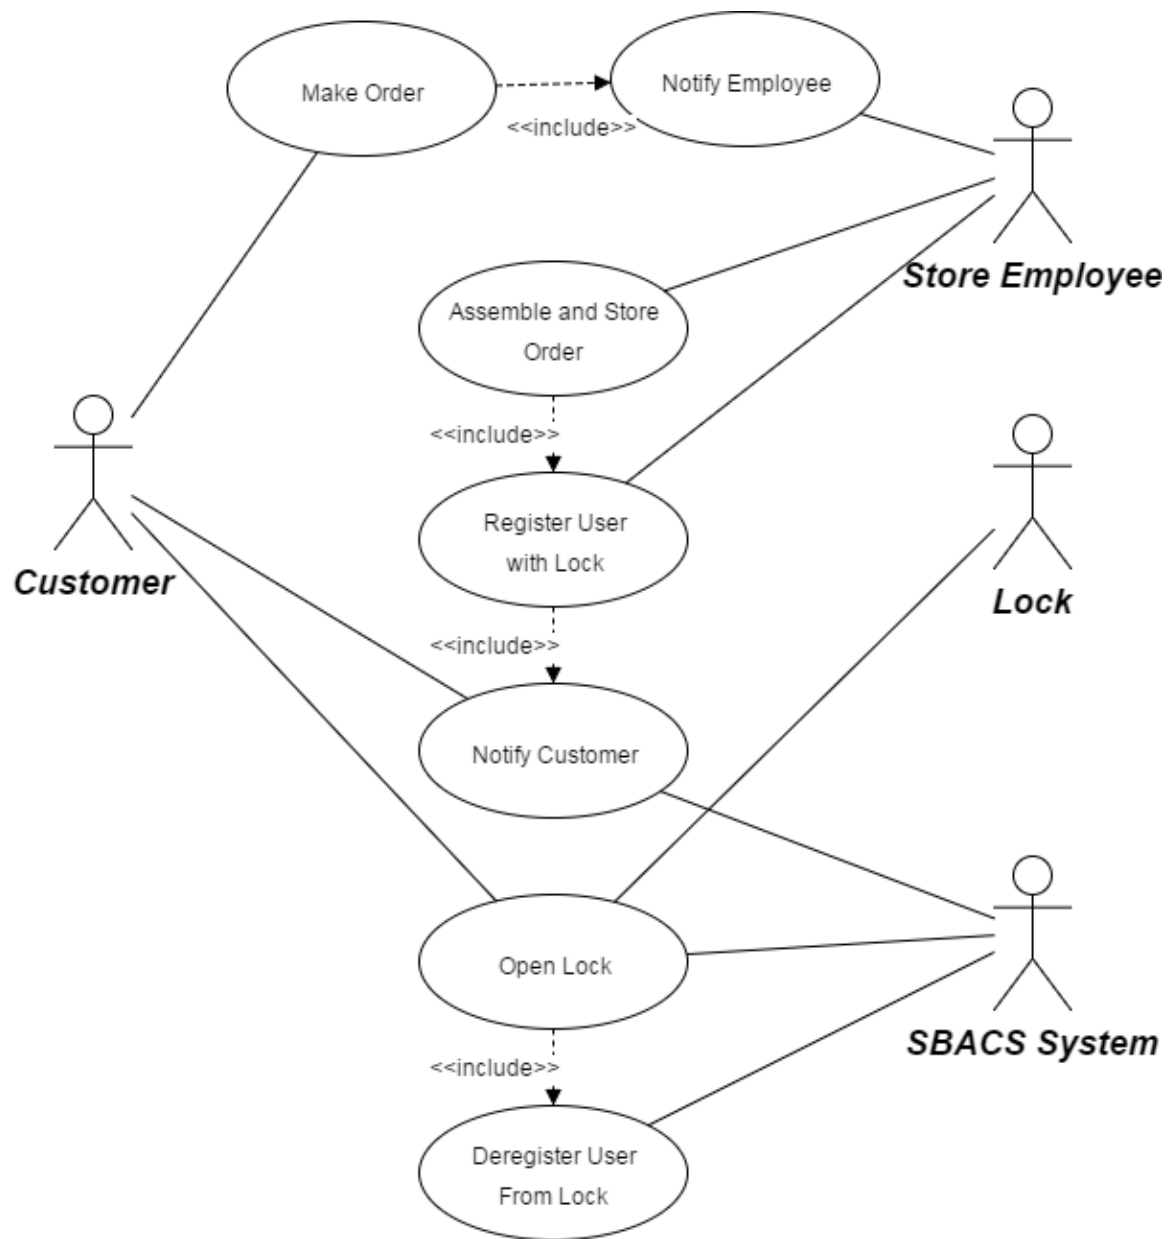


Figure 1: Online Order Secure Pickup Use Case Diagram

4.2 Central Mail Package Pickup

4.3 Long Term Storage

4.4 Service Provider

Chapter 5

Problem Analysis and System Design

5.1 Overall System Analysis

The requirements of our system from the use cases described above lead us to a simple outline for the entire system. The main component of the system would be the cloud server, which would store information regarding users, locks, and the relationships between them. The lock hardware would need to be able to securely accept data from the user, and securely send that data to the server for verification. Since handling electronic security can be difficult to do manually, having an application that handles much of the busy work would greatly help users interact with the system.

One of the main goals of the system was to improve the reusability of the electronic keys that are created. For this reason, the idea of a key was split into two parts: an authenticator, and an identity. Users would be able to create many authenticators, which were things like a password, or a PIN. These then could be combined into an identity, which would be used to open a particular lock. This allowed making identities that reuse the same authenticator, and reusing identities for different locks.

Another goal of the system was to provide additional behaviour on top of just being able to unlock particular locks. This extra behaviour was captured in the notion of a registration, which ties one lock to one identity. Through the registration,

we were able to implement notifications to the user as well as streaming of the lock's contents.

5.2 NFC

Determining how NFC communication should take place required analysis of two hardware systems: NFC card readers, as well as mobile devices. The most desirable protocol would be able to handle the widest variety of available hardwares for the two devices. The performance considerations between modes was fairly minimal, so preference was placed on the portability of the solution.

5.2.1 Card Readers

Since NFC cards are primarily designed for NFC communications, there were few restrictions that stemmed from potential choices in card reader. Since NFC communications are specified by the ISO, most cards support enough protocols that any decision on our part would be very likely to be supported by any card that would be desirable for any other reason.

5.2.2 Mobile Devices

The two most popular operating systems for mobile devices are iOS and Android []. Since iOS devices have NFC disabled for everything except Apple Pay[], the only option that remained was Android. Apple devices would represent a large part of the potential market, so alternatives to NFC would have to be considered.

Among Android devices, there exist devices which have hardware support for NFC communications, and those which rely on host-based card emulation. Devices with hardware support have a component called a Secure Element which performs all of the communication with the external NFC terminal. Later, applications can query this element to determine the status of the transaction, as well as other data. Devices which use host-based card emulation use a software implementation of secure elements. Since host-based card emulation is done through software, it will run on

all Android devices running version 4.4 or greater[], which represents over 99% of all devices currently in use.

Android offers an API called Beam which is the only way Android devices can use NFC in active mode []. Beam, however, does not support sending more than one message between devices. Since the information we are sending can be fairly large in the interest of security, this was not feasible given the restrictions of the NFC protocols we used. Further, active communications are easier to eavesdrop on, as discussed in the background section. We decided that these costs outweighed the simplicity of the Beam API, so passive communications were chosen for the implementation.

Since more than one application could potentially want to handle an NFC message, the message protocol contains a field called the Application IDentifier, or AID []. These AIDs are just large numbers which identify which application should handle the associated message. The Android operating system is responsible for delivering the APDUs to the appropriate application's service []. Reserving a particular AID costs money, so we decided that the probability of a collision occurring when taking a public AID would be an acceptable cost for the purpose of our demonstrative implementation.

5.2.3 Protocol

Since our NFC communications may require more data than can be fit within an Application Protocol Data Units (APDUs), we required a protocol which would handle segmenting and recombining the message. APDUs are defined in ISO 7816-4 [] and are the units used by ISO 14443-4 [], which describes the transmission protocol used by NFC devices. They are restricted to 256 bytes, including headers.

To work around this, the hardware device connected to the shield maintains a buffer. Under ISO 14443-4, messages can be reliably transferred, so managing this buffer is the main consideration of our protocol. The hardware determines the maximum amount of data that can be stored in one APDU, and fills in this value into the length field of the APDU that it sends to the Android device. Then, the Android application responds with the minimum of that much data, and all of the remaining

data that it has to send. Once the hardware receives an amount of data less than the potential maximum, it deactivates the connection. In the event that the data from the Android application fits exactly into the last message that would be sent, the protocol still works, as the application will then respond with zero data bytes.

5.3 Android

The Android application's goal was to provide the end users with an easy way of interacting with the SBACS system. Therefore the application was designed to be as simple as possible while still maintaining flexibility when dealing with the cloud server as well as the many potential hardware devices.

5.3.1 User Interface

The natural decision for designing the basic workings of the application was to use the Model-View-Controller (MVC) design pattern. This pattern separates the underlying data (model) from the display that the user sees (view) as well as the components that the user interacts with (controller). Android provides APIs to support MVC. Model objects can be simple objects, but view objects can subclass the View class or its more specific subclasses, and the controller objects can subclass the Adapter class or its subclasses.

Android applications themselves should follow the patterns set out by the standards. Activity classes represent the pages that contain the various views and controllers that users can interact with. We decided on a design where the user first encounters a login Activity, which prompts them to either sign up or log in. Once the user has logged in, we provide a hub Activity which leads to the various other Activities in the application. These other Activities show the various data related to the user, such as their authenticators and identities.

Communication between Activities is handled by a class provided by Android called an Intent. Intent objects contain information about the nature of the request to begin the new Activity. In this way, information such as the user's identification

number could be sent from the login Activity to the hub Activity, which would allow the application to correctly load the information from that user when displaying the data Activities.

5.3.2 Notifications

The design for the notifications that would inform end users of information such as their newly available registration with a particular lock was based partly on decreasing load on the server. Having the server maintain a large amount of session information as well as spending time and memory on the timers to be able to update the user of these changes was thought to be too great a cost. For this reason, it seemed simpler and more effective to have the application poll the server at a particular access point for new information.

This polling was implemented simply in the application using a service which runs in the background of the Android device. This service regularly hits the server at a particular endpoint designed for handling these notifications. Since many users may be using the application at once, it was important to consider the performance of the endpoint's code. The endpoint returns information valuable to the application for display in the notification. The notifications make use of Android notification's ability to launch an Activity with an Intent to take the user to the appropriate Activity associated with the notification.

5.3.3 NFC Handler

As mentioned earlier, the NFC communication method that we decided on used host-based card emulation. Android provides an API to accomplish these communications, through the `HostApduService` class [1]. This class provides a framework for creating a service which runs in the background on the Android device. The service handles setting up and closing NFC communications, while also providing overrideable methods to determine what data should be sent based on the incoming message.

Services also take an Intent object to determine their purpose. This feature was used to have the login Activity start the service with the correct key information for

a user's NFC authenticator. This way, the data sent could be overridden with the user's secret key value. Since passive NFC communications are difficult to eavesdrop on, this was thought to be a safe method of conveying the secret key to the lock without requiring too much data to be sent over the slow NFC physical links. This passing of information of course was also configured to follow the protocol designed above for NFC communications.

5.4 Lock Hardware

Based

5.4.1 Core Circuit

Design of the system core began by laying out the requirements for the core:

- There should be LEDs to indicate the status of the system, and to notify the user of the results of their access attempt
- There must be an electrically-powered lock, and a means to digitally alter its status
- The core must run a process where all authentication data will be passed, and this process must communicate to the server over the internet
- The core must have USB ports, to support having modules send authentication data to the core

Figure 2 is a block diagram of the first draft of the core layout.

5.4.2 Core Software

5.4.3 Authentication Modules

To meet the requirement of having a modular framework that supports multiple authentication methods, modules that perform the gathering of authentication data




Figure 2: Hardware Block Diagram

can be attached to the system, to increase the authentication options. As a result, multiple design patterns emerged immediately in the design of the software for the modules. Interfacing with the system core would be the same across all authentication modules, so only the method used to gather authentication data is unique to each module. Two design patterns can be extracted from this statement: an Adapter pattern, where the Adaptee provides the interfaces necessary for communicating with the core, and a Template method for handling the authentication data. The Adapter can simply implement the Template method, as well as add any other functionality necessary for the means of authentication.

Talk about actors-boundaries-controllers-entities here

In the planning stages of the project, it was intended that the authentication modules would communicate with a central single-board microcontroller, such as a Arduino, which would be deemed a "communication board." This communication board would be responsible for receiving data from the authentication modules, and for error detection and error correction, removing that overhead from the central

computer. Once all of the authentication data was gathered, the board would send the data over a USB connection to the central computer. However, this design was tossed, for multiple reasons:

1. The Arduino, as well as most other microcontrollers in the hobby market, can only communicate with another Arduino over I2C;
2. There is only one pair of pins reserved for I2C on most microcontrollers, meaning there can only be one "slave" for each "master" communication board (refer to pinout here);
3. Connecting all of the authentication modules over USB allowed for far more modularity than the "communication board" would have.

Maybe add a figure of the Arduino pinout here?

Thus, the idea of using I2C communication was scrapped, in favour of simply connecting each module to the core over USB. This simplified the development of communication modules considerably.

5.5 Cloud

5.5.1 Amazon Web Services

Amazon Web Services (AWS) provides a platform to host various cloud services, including web servers and data storage options [2]. In this project, Amazon Web Services was chosen to host the web server and the database. This decision was made because the platform is versatile and easy to use, providing a great level of control where needed but also handling the complicated configuration details automatically. AWS also provides these services for free within a trial period of one year, making it an even more attractive choice for a project with limited lifetime. To host the web server, the choice of programming language and technology, as well as the choice of operating system to run the server, are left to the user. Similarly, the database technology is a choice of the user, making AWS a very flexible platform to work with.

Amazon Web Services also provides some other very helpful features, such as domain name services, SSL certificate management, and various security configuration details to protect the web server and our data. We chose to use AWS because of the variety of services that were offered that provide everything we needed to effectively host our web server for this project.

5.5.2 Representational State Transfer

The Representational State Transfer (REST) architectural style [3] was chosen to be used for our web server. The main principle of REST that we adhered to is the idea of a stateless server in the client-server relationship. This means that the server does not store any session information; enough session information must be passed from the client on each request to identify the necessary information for the request. The benefit of this for the server is that less information is stored on the server, so less time will be spent contacting the database to determine session information. Another feature of REST is the idea of a uniform interface. This helps to keep the system simple by relying on identifiers to specify which resource or service is needed at the time. The path and the query string in the URL are examples of identifiers that specify which endpoint or service is being requested. The web server for this project has this uniform interface as all HTTP requests go to a single Front Controller, which then determines based on the identifiers the appropriate handler of the request.

5.5.3 JavaScript Object Notation

In order to send information between the client and the server, a standard format for representing the data is necessary. We have chosen to use JavaScript Object Notation (JSON) to serve this purpose. "JSON is a lightweight, text-based, language-independent data interchange format" [4]. JSON provides an easy way of transferring data that is flexible and has support from many programming languages, making it a very good choice for our project, where different components will be using different technologies and built in different languages. It is also very lightweight in terms of its grammar, having very few elements that allow for many different types of data,

keeping the parsing easy which can help to improve response time.

5.5.4 Node.js

As mentioned above, Amazon Web Services allows for web servers written in many languages, leaving the choice of language to the user. In our case we decided to use Node.js to write our server code. This language was chosen because it is able to run on multiple operating systems and does not require many tools for effective development. The ability to run on multiple operating systems is very useful to test our code in different ways, and leaves us with flexibility of which operating system to host the server with. In the early stages of the project, we considered using C# with the .NET framework to build the web server, as some members of the group had experience working with this particular technology in the past for a similar purpose. After some discussion, it was decided that in order to effectively develop the application in C# tools would be needed that are not available on the computers at school, which would limit our ability to work on the server. This is a major reason that we decided on Node.js, since it can be very effectively developed using only simple text editors, then can be executed quickly, without needing to install any build tools. This portability made it very easy to work on the code in the same development environment whether we were at school or at home.

5.5.5 MySQL

To store data that is needed by our system, we decided to use a relational database. All members of the group have taken at least one course that has talked about relational databases, so we all had at least some basic understanding of this type of database. MySQL was chosen because it was known to members of the group, making it easy for everyone to make queries to the database whenever it was needed.

5.5.6 Relational Database Theory

We follow third normal form in our database.

5.5.7 Authentication for Server API

Many requests to the server allow users to read, modify, and delete data, so we needed a way to ensure that users were identified and authenticated to allow certain requests to be performed. We used Hashed Message Authentication Codes (HMAC) to accomplish this goal. The user is able to login using their username and password, and upon successful login, they are provided with a key and user id. The key is then used to hash the body of subsequent messages. The client must also send the user id to identify who is trying to make the request. When the server receives the request, it will hash the body of the message as well using the key that it has stored for that user. If the hashed body that is provided is the same as what the server gets from hashing the body, then this means that the user is authenticated. The provided user id is used to determine which key the server will use to hash the message, and is also used to determine whether or not the user is authorized to make the request they are attempting to make.

5.5.8 Storing Authenticator Data Securely

In order to protect user passwords from attacks, the database must not store the passwords in plain text. If the passwords are in plain text, an attacker could simply read the data and know the users password. To prevent this problem, we use salting, which allows us to store a hashed version of the data, along with a salt that is used to generate the hashed value. To compare a given password by the user to the value that is stored, the server must retrieve the salt and use it to hash the given password. The hashed value in the database should then match the hashed given password, meaning that the user has provided the same password. With this technique, the data is never stored in plain text so it is less vulnerable to many types of attacks.

5.5.9 Sending Data Securely Over the Network

HTTP is not a secure protocol. HTTPS is secure.

5.6 Lock Demonstration

This chapter will be about the design and construction of the demo unit used at the poster fair

Chapter 6

System Implementation

6.1 Android

The Android application was developed using Android Studio exclusively, and is primarily composed of Java classes, XML documents, and gradle build files. Android studio was selected since it seemed to be the natural tool for writing an Android application, and didn't have any negatives with respect to our requirements. Gradle was chosen as the build tool as it is naturally integrated with Android studio.

6.1.1 Server Interfacing

Connections to the server were managed through the android Volley API. Volley provides a simple interface over which HTTP and HTTPS messages can be sent. Standard operation of Volley follows a simple procedure: first a RequestQueue is set up, and then various Requests are enqueued. These requests are then dequeued by the RequestQueue's thread, which then creates an HTTP/HTTPS connection. The response that comes back over the connection is handled by another thread.

The code for handling these responses is put in a class that extends a Response.Listener of a given type. Since these classes have just a few relatively simple methods, anonymous classes were used in all cases. Further, we decided that handling the parsing should be done as safely as possible, so all responses were first parsed simply as

Strings. Then, the response was attempted to be parsed as a JSON object of the type expected from the server. If that parsing failed, the message was instead considered a error, and handled from there.

6.1.2 HMAC

Since the server used HMAC to handle security issues, the application needed to make frequent use of the headers that the server expected to handle authentication. For this reason, a helper class called HMACHelper, was created. This class provided methods that were commonly used by all requests about data particular to the current user. Most notably, HMACHelper provided a method which calculated the secret using the same algorithm that the server would use. To reiterate how HMAC functions, the body of the message is hashed using a private key shared by the server with the application. When the application sends data that requires authentication, the server checks that the secret value provided by the application matches what it calculates using the body of the message as well as the private key associated with the user. These private keys are generated by the server when the user logs in to the application, and expire over time, creating the notion of a session.

The precise algorithm used for the hashing function matches the decision made for the server. This is necessary, since otherwise the secret values would not match between the application and the server. The application relies on a version of the Password-Based Key Derivation Function algorithm using SHA256 to be available on the Android device. The only workaround for this added requirement would have been to implement the algorithm (specified in RFC 2898 with test vectors in RFC 6070) ourselves, however doing so is known to be quite dangerous as any minor error could result in a security breach. In addition, the number of iterations and the resulting length also had to be kept as the same value.

6.1.3 Entity Representation

The application needed to be able to display data to the user that came from the server and database. These data are represented in the MVC design of the application.

Since the database is relational, the model classes that represent the constructs stored in the database, such as identities, authenticators, locks, and registrations, were all passive entity classes whose only real role was to contain data.

Each kind of data that was of interest to the user was given its own Activity, which is similar to the notion of a page in a web application. These Activities were all primarily composed of a ListView object, which displays a list of data. This list was configured with an Adapter class which extends the general `ArrayAdapter<T>` class. These adapters were used to convert the model classes into a view. The adapter classes were composed of an XML file which described the layout of the view, as well as a java class which provided the methods by which the data could be used to fill in the layout.

6.2 Hardware

The hardware portion of the project consisted of two major portions: the

Since driving the LEDs and the electrical lock only requires digital output, no analog pins would be necessary. While the single-board microcontrollers considered had optional boards which could add Wifi or Ethernet capabilities, they would not have been a sufficient solution to run the core process, as they were lacking a sufficient number of USB ports, and the overhead for implementing the core would have been too high. Because the microcontrollers do not run an operating system,

A Raspberry Pi was ultimately selected as the computer to run the core process. The Pi 3 Model B has four USB ports, built-in Wifi capability, and forty pins for general-purpose IO (GPIO). These hardware features made the Raspberry Pi the best out-of-the-box solution to run the core process.

6.2.1 Lock Control Circuit

The next design problem to consider was the means to control the status of the electric lock. The simplest locks available, solenoid locks and electric strike locks, do not include any control mechanisms. They are either powered on, to disengage the

lock, or powered off and locked. Therefore, it was necessary to implement a circuit to control the power to the lock, with the switching being driven by one of the Raspberry Pi's GPIO pins. There were two options considered for this:

1- Transistor/diode circuit

This circuit would use a TIP120 Darlington transistor, coupled with a 1N4004 flyback diode, to drive the lock. The circuit is demonstrated in figure NEXT.

Figure of circuit

Some heat-related math for the transistor: The power law for a single voltage drop is $P = IV = I^2R$. A transistor can be treated as two voltage drops, one from base to emitter (V_{BE}), and another from collector to emitter (V_{CE}), resulting in the power dissipated by a transistor being given by:

$$P = V_{BE}I_B + V_{CE}I_C$$

Since $I_B \ll I_C$ when the lock circuit is active, the $V_{BE}I_B$ term is negligible in this case, so the power dissipated by the transistor is given by:

$$P = V_{CE}I_C$$

At steady state, the voltage drop from the solenoid is negligible, as $\frac{di}{dt}$ is zero, resulting in $V = L\frac{di}{dt} = 0$. This means that the voltage drop across the transistor, from emitter to collector, is 12V. Given that the solenoid draws 650 mA of current, the power across the transistor is:

$$P = 12\text{ V} \times 650\text{ mA}$$

$$P = 7.8\text{ W}$$

This power dissipation is quite significant in comparison to, for example, the resistors in the LED circuit, which at most will dissipate about 50 mW. (n.b. should I put the math behind this in an appendix?) A heat dissipation mechanism would be necessary to ensure that the system does not heat up dangerously while in operation, which may not be practical in a real-life scenario. Secure systems should be completely enclosed, as a fan would introduce a point of physical access to the system.




Figure 3: Relay Circuit Diagram [5]

Insert some more math and talk about the flyback diode to deal with the voltage spike - take the lock to a DOE lab at some point and measure the voltage spike when 12V DC is applied, to see what the minimum RB breakdown voltage should be

2- Relay circuit

Internally, a relay is very similar to the transistor/diode circuit, with a mechanical switch added. The relay itself has three inputs, and three outputs. The inputs are a positive voltage (V_{CC}), ground (GND), and an input signal (IN). The outputs are a normally-closed terminal (NC), a common terminal (CO), and a normally-open terminal (NO). The circuit diagram of a relay is shown in Figure 3.

What separates the relay from a transistor/diode circuit is that it uses a solenoid internally to flip a mechanical switch. When the IN signal is low (0 volts), the switch short-circuits the "normally closed" connection, allowing current to flow from the NC terminal to the CO terminal. The NO terminal is open-circuited, so no current will flow. When the IN signal is high (5 volts), the solenoid inside the relay generates a magnetic field, which causes the switch to instead short-circuit the "normally open" connection, allowing current to flow from the NO terminal to the CO terminal, while the NC terminal is open-circuited.

6.2.2 Lock Control Software

6.2.3 Authentication Modules

6.3 Cloud

Chapter 7

Testing and Bug Fixes

Chapter 8

Conclusions

Appendix A

Sample Appendix

Bibliography

- [1] Tony R. Kuphaldt. *Chapter 3 - Electrical Safety*. URL: <https://www.allaboutcircuits.com/textbook/direct-current/chpt-3/ohms-law-again/> (visited on 03/07/2017).
- [2] Amazon. *What is AWS? - Amazon Web Services*. URL: <https://aws.amazon.com/what-is-aws/> (visited on 03/09/2017).
- [3] Roy T. Fielding. *Fielding Dissertation: CHAPTER 5: Representational State Transfer (REST)*. URL: http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm (visited on 03/09/2017).
- [4] ECMA International. *Final draft of the TC39 The JSON Data Interchange Format standard*. URL: <http://www.ecma-international.org/publications/files/ECMA-ST/ECMA-404.pdf> (visited on 03/09/2017).
- [5] LC Circuits. *Relay Module*. URL: <http://elecfreaks.com/store/download/datasheet/breakout/Relay/Relay.pdf> (visited on 10/19/2016).